

Assignment 2:

- 1) Following are the results of running an SVM implemented with scikit on the four provided datasets. A linear, polynomial($d=2$), and RBF kernel were tested on each dataset.

Data Set: twospirals.dat

Support Vector Machine(Kernel = linear):

PPV: 0.65053 NPV: 0.63619 Specificity: 0.66800

Sensitivity: 0.61800 Accuracy: 0.64300

Support Vector Machine(Kernel = poly):

PPV: 0.61966 NPV: 0.53655 Specificity: 0.82200

Sensitivity: 0.29000 Accuracy: 0.55600

Support Vector Machine(Kernel = rbf):

PPV: 0.95219 NPV: 0.95582 Specificity: 0.95200

Sensitivity: 0.95600 Accuracy: 0.95400

Data Set: twogaussians.dat

Support Vector Machine(Kernel = linear):

PPV: 0.97087 NPV: 0.96907 Specificity: 0.96907

Sensitivity: 0.97087 Accuracy: 0.97000

Support Vector Machine(Kernel = poly):

PPV: 0.97073 NPV: 0.96410 Specificity: 0.96907

Sensitivity: 0.96602 Accuracy: 0.96750

Support Vector Machine(Kernel = rbf):

PPV: 0.98522 NPV: 0.96954 Specificity: 0.98454

Sensitivity: 0.97087 Accuracy: 0.97750

Data Set: clusterincluster.dat

Support Vector Machine(Kernel = linear):

PPV: 0.41657 NPV: 0.00000 Specificity: 0.00000

Sensitivity: 0.71400 Accuracy: 0.35700

Support Vector Machine(Kernel = poly):

PPV: 1.00000 NPV: 1.00000 Specificity: 1.00000

Sensitivity: 1.00000 Accuracy: 1.00000

Support Vector Machine(Kernel = rbf):

PPV: 0.99602 NPV: 1.00000 Specificity: 0.99600

Sensitivity: 1.00000 Accuracy: 0.99800

Data Set: halfkernel.dat

Support Vector Machine(Kernel = linear):

PPV: 0.68346 NPV: 0.81918 Specificity: 0.59800

Sensitivity: 0.86800 Accuracy: 0.73300

Support Vector Machine(Kernel = poly):

PPV: 0.70721 NPV: 1.00000 Specificity: 0.58600
Sensitivity: 1.00000 Accuracy: 0.79300
Support Vector Machine(Kernel = rbf):
PPV: 1.00000 NPV: 0.99800 Specificity: 1.00000
Sensitivity: 0.99800 Accuracy: 0.99900

2) Performance Notes:

Two Spirals dataset:

For this dataset, as expected, the linear classifier performed poorly with 64% accuracy however the polynomial performed even worse with a 56% accuracy. This is because the data is not linearly separable so the linear performed poorly, and additionally in the polynomial space, the data must not be linear separable either. However, perhaps adding more features with a degree 3 polynomial kernel would improve the classification. Finally, the RBF kernel performed exceptionally well, with an accuracy of 95%. This is because the RBF kernel effectively measures the Euclidean distance between two points, so here, it acts similar to the way a k-NN neighbour classifier would act and since the classes are in two distinct regions, judging a sample's class by its nearest neighbours is an effective classification strategy.

Two Gaussians dataset:

For this dataset all three performed equally as well with an accuracy of $97\% \pm 0.75\%$. This is because the data is already linearly separable in the data's initial feature space, so the linear kernel works well and the other two are just overkill.

Cluster in Cluster dataset:

For this dataset, the linear classifier performed poorly with 35% accuracy (worse than random). However both the polynomial and RBF kernel achieved a $99.9\% \pm 0.01\%$ accuracy. This is because the polynomial features added by the polynomial function were able to be linearly separated in the new feature space. Additionally, the RBF kernel was effective due to the data lying in distinct regions.

Half Kernel dataset:

For this dataset the linear classifier performed better than expected with a 73% accuracy however upon inspection of the plotted classified data, it is clear the SVM just found the optimal place to draw a straight line for best classification results. The polynomial performed mildly well with a 79% accuracy but once again, inspecting the plotted results shows the classifier is misclassifying a large region and perhaps a degree 3 kernel would perform better. Finally, the RBF kernel performed exceptionally well with a 99.9% accuracy. This performance is due to the same reason as before, the fact that the data is in two distinct regions.

3) Following are the ROC curves obtained for the RBF classifier on each dataset, along with the points that were used to construct the ROC curve. The x-axis and y-axis are extended by 0.025 as some points would be lying directly on the edges of the graph if the two axes went from 0 to 1.


Note: each of the datasets had a parameter(C, γ) pair that resulted in a classifier with both minimum false positive rate and maximum true positive rate, meaning that when the convex hull was taken to remove the sub-optimal classifier points, only one remained, creating a square ROC curve. i.e.) Because all the removed points all had more false positives and less true positives than the point that remained

in the convex hull, regardless of which efficiency measure we prioritize, the same $(C, \gamma\gamma)$ pair remained the “best”, so a single $(C, \gamma\gamma)$ pair defined the shape of the ROC curve.

Note II: Each Data set was tested with all combinations of C and $\gamma\gamma$ values where $C = 2^{ii}$


and $\gamma\gamma = 2^{jj}$ such that $ii \in [0,20]$ and $jj \in [0,8]$ giving a total of 189 points, except for the Two Spirals dataset, which was tested with $ii \in [-20,20]$ and $jj \in [-8,8]$, yielding 697 points, as a totally optimal solution was not found in the initial smaller range.

Two Spirals dataset:




Here, the SVM's efficiency was a lot more dependent on its parameter's values with some parameters as bad as a random classifier and the best parameters being $C = 2^{-2}$ and $G = 2^1$ achieving a TRP of 96.2% and an FPR of 3.8%

Two Gaussians dataset:




Here, all 189 points are above 90% TPR and below 15% FPR

Cluster in Cluster dataset:



Here, all 189 points are at or near 100% TPR and 0% FPR

Half Kernel dataset:




Here also, all 189 points are at or near 100% TPR and 0% FPR

- 4) The SVM-R was able to improve upon the K-NN neighbours by 1-4% accuracy depending on the data set. This similar performance is due to the fact that k-NN neighbours and RBF kernel both use a similarity function(Euclidian distance) between two samples as features.


The RBF kernel with Euclidian distance highlighted in red:
$$K(x_i, x_j) = \exp \left(-\frac{\|x_i - x_j\|^2}{2\sigma^2} \right)$$

Since the data lies in two distinct regions, both classifiers were able to accurately classify the data. The increase in performance, however small, could be due to the RBF's ability to weight certain neighbours as less important, so a sample on the edge of the two regions, that is very close to both positive and negative samples would still be closer to the samples lying in the centre of one region than the samples lying in the centre of another and thus can be more accurately classified if we ignored the samples that are near but also lying on the edge of the regions. This is illustrated below.

Dark Blue = positive sample, **Red** = negative sample, **Black** line = "actual" class boundary.



2 of the 3 nearest neighbours are positive so classify as positive.



Let the **orange** crosses denote highly weighted negative samples, and let the **light blue** circles denote highly weighted positive samples. Here we can see the sample to be classified is close (more similar) to the **orange** crosses and so we'd expect this sample to be classified as negative

Implementation Notes:

The SVM's were implemented using scikit's sklearn library. For implementing the ROC curve, a self implemented solution was used rather than using a library. The ROC curve was obtained by first running the SVM-R with every combination of parameter values in a certain range and saving their respective true positive and false positive rates. Next three points were added to the set of (TPR, FPR) pairs to be plotted; (min(FPR), 0.0), (1.0, max(TPR)), and (1,0). To illustrate the reason for adding these points, refer to the image below. Next, the convex hull was found using scipy's `ConvexHull()` function and the points in the hull were plotted to obtain the ROC. To obtain the AUC, numpy's `trapz()` function was used which uses the composite trapezoidal rule to calculate the area under a set of points with straight lines drawn between them. This is the same method the sklearn's `auc()` function uses.

